

Introduction to inverse reinforcement learning (2/3)


Eiji Uchibe

Dept. of Brain Robot Interface

ATR Computational Neuroscience Labs.

Naïve Imitation Learning Works Poorly

- The robot's errors compound when drifting away from the expert's demonstrated trajectories




- No mechanism to recover when the generated trajectory deviates from the demonstrated ones

Preparation

- $\tau = \{s_1, a_1, \dots, s_T, a_T\}$: trajectory of state s and action a (green)
- $r(s, a; w)$: immediate reward at state s and action a parameterized by w
 - For a linear function approximator,
 $r(s, a; w) = w^\top \phi(s, a)$,
where $\phi(s, a)$ is a feature vector
- Total reward along a trajectory τ

$$R(\tau) = \sum_t r(s_t, a_t; w) = w^\top \sum_t \phi(s_t, a_t) = w^\top \phi(\tau)$$



Kalakrishnan, M., Pastor, P., Righetti, L., & Schaal, S. (2013). [Learning objective functions for manipulation](#). *Proc. of ICRA*, 1331–1336.


Maximum Entropy Inverse RL

- $P(\tau|w)$: probability of τ parameterized by w (blue)
- The expectation of features under $P(\tau|w)$ is equal to the expected empirical feature count

$$\underbrace{\sum_{\tau} P(\tau | w) \phi(\tau)}_{\text{expected feature under } P(\tau | w)} = \underbrace{\frac{1}{m} \sum_j \phi(\tau_j)}_{\text{expected empirical feature count of the expert}} \triangleq \tilde{\phi}$$

expected feature under $P(\tau | w)$ expected empirical feature count of the expert

- Since many reward weights satisfy the above constraint, the principle of maximum entropy is employed to resolve ambiguities in choosing distributions



Principle of maximum entropy (1/2)

- Choose a 'best' from a number of different probability distributions that all express the current state of knowledge
- Constrained optimization problem

$$\max_p - \sum_{\tau} p(\tau) \ln p(\tau) \quad \Rightarrow \quad \text{Maximize the entropy of } p(\tau)$$

$$\sum_{\tau} p(\tau) \phi_i(\tau) = \tilde{\phi}_i, \forall i \quad \Rightarrow \quad \text{constraints on the expected feature}$$

$$\sum_{\tau} p(\tau) = 1, \quad p(\tau) \geq 0 \quad \Rightarrow \quad p(\tau) \text{ is a probability distribution}$$

Principle of maximum entropy (2/2)

- The Lagrangian is given by

$$\mathcal{L} = - \sum_{\tau} p(\tau) \ln p(\tau) + \lambda \left(\sum_{\tau} p(\tau) - 1 \right) + \sum_i w_i \left(\sum_{\tau} p(\tau) \phi_i(\tau) - \tilde{\phi}_i \right)$$

– λ, w_i : Lagrange multipliers

- Calculus of variation: $\partial \mathcal{L} / \partial p = 0$ yields

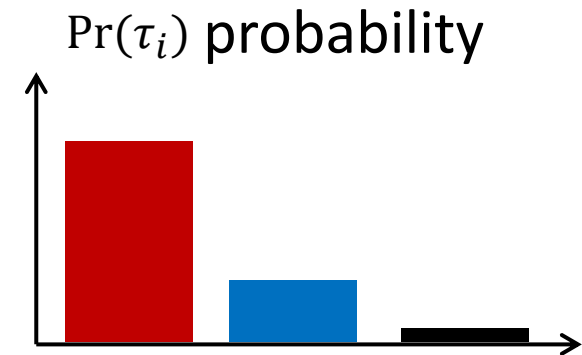
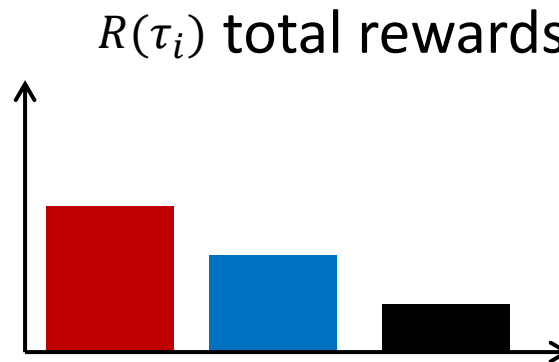

$$\ln p(\tau) = \sum_i w_i \phi_i(\tau) + \lambda - 1$$

$$\Rightarrow p(\tau|w) = \frac{1}{Z(w)} \exp(w^\top \phi(\tau)) = \frac{1}{Z(w)} \exp(R(\tau))$$

Probability under the principle of maximum entropy

- Trajectories with equivalent returns have equal probabilities
- Trajectories with higher returns are exponentially more preferred

$$p(\tau|w) = \frac{1}{Z(w)} \exp(R(\tau; w))$$



Parameter optimization

- w is estimated by maximum likelihood from the expert's dataset

$$\mathcal{D}^\pi = \{\tau_j^\pi\}_{j=1}^N$$

- Log-likelihood is given by

$$L(w) = \frac{1}{N} \sum_{\tau_j^\pi \in \mathcal{D}^\pi} \ln p(\tau_j^\pi \mid w) = \frac{1}{N} \sum_{\tau_j^\pi \in \mathcal{D}^\pi} R(\tau_j^\pi; w) - \ln Z(w)$$

- w is updated by gradient ascent

$$w \leftarrow w + \alpha \nabla L(w), \quad \nabla L(w) = \frac{1}{N} \sum_{\tau_j^\pi \in \mathcal{D}^\pi} \nabla R(\tau_j^\pi; w) - \nabla \ln Z(w)$$

- How to evaluate the gradient of $\ln Z(w)$?

MaxEnt-IRL: evaluate $Z(w)$ by solving model-based reinforcement learning

- $$\begin{aligned}\nabla \ln Z(w) &= \frac{1}{Z(w)} \frac{\partial Z(w)}{\partial w} = \frac{1}{Z(w)} \sum_{\tau} \exp(w^{\top} \phi(\tau)) \phi(\tau) \\ &= \underbrace{\sum_{\tau} P(\tau; w) \phi(\tau)}_{\text{expected feature under } P(\tau; w)} = \mathbb{E}_{P(\tau; w)}[\phi(\tau)]\end{aligned}$$

- In RL, the following equation holds

$$\sum_{\tau} P(\tau; w) \phi(\tau) = \underbrace{\sum_{s, a} P(s, a; w) \phi(s, a)}_{\text{expected feature under the joint state-action distribution}}$$

Samples from $P(s, a; w)$ can be generated by the optimal policy $\pi^*(a | s; w)$ trained with the current estimated reward $r(s, a; w)$

expected feature under the joint state-action distribution

Simulation: FrozenLake task

- Discrete-state and discrete-action MDP task provided by OpenAI gym
 - Simplify the task by considering deterministic MDP
- Positive reward (+1) for entering the goal 📦 from left
- 0 otherwise
- Episode ends when entering a hole 🕒
 - 0 reward is given to the agent
- <https://uchibe.github.io/BILT-B/>

